

# **EE4077 Fundamentals of Machine Learning**

## Linear Regression – I

---

Fall 2021

EE–Marmara University

## 1. Recap of MLE/MAP

## 2. Linear Regression

Motivation

Algorithm

Univariate solution

Multivariate Solution

Probabilistic interpretation

Computational and numerical optimization

# Linear Regression

---

## Recap of MLE/MAP

### Linear Regression

- Motivation

- Algorithm

- Univariate solution


- Multivariate Solution

- Probabilistic interpretation

- Computational and numerical optimization

# Task 1: Regression

*How much should you sell your house for?*



**input:** houses & features   **learn:**  $x \rightarrow y$  relationship   **predict:**  $y$  (*continuous*)

**Course Covers:** Linear/Ridge Regression, Loss Function, SGD, Feature Scaling, Regularization, Cross Validation

# Supervised Learning

## Supervised learning

In a supervised learning problem, you have access to input variables ( $X$ ) and outputs ( $Y$ ), and the goal is to predict an output given an input

- Examples:
  - **Housing prices (Regression)**: predict the price of a house based on features (size, location, etc)
  - **Cat vs. Dog (Classification)**: predict whether a picture is of a cat or a dog

## **Predicting a continuous outcome variable:**

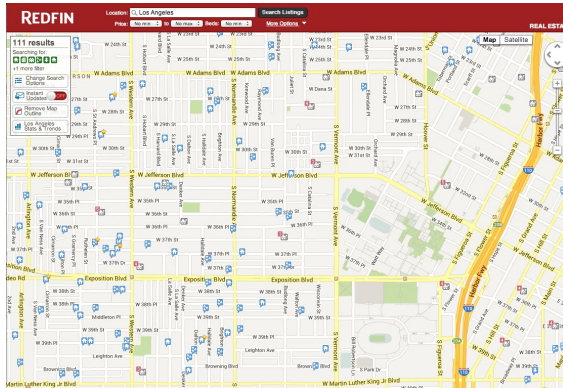
- Predicting a company's future stock price using its profit and other financial info
- Predicting annual rainfall based on local flora and fauna
- Predicting distance from a traffic light using LIDAR measurements

## **Magnitude of the error matters:**

- We can measure 'closeness' of prediction and labels, leading to different ways to evaluate prediction errors.
  - Predicting stock price: better to be off by 1\$ than by 20\$
  - Predicting distance from a traffic light: better to be off 1 m than by 10 m
- We should choose learning models and algorithms accordingly.


## Ex: predicting the sale price of a house

Retrieve historical sales records  
(This will be our training data)





# Features used to predict



**3620 South BUDLONG**  
Los Angeles, CA 90007  
Status: Closed

**\$1,510,000**  
Last Sold Price


**14**  
Beds

**6**  
Baths

**4,418** Sq. Ft.  
6342 / 86, Ft.

Built: 1956 Lot Size: 5,548 Sq. Ft. Sold On: JUL 26, 2013

[Overview](#)
[Property Details](#)
[Tour Insights](#)
[Property History](#)
[Public Records](#)
[Activity](#)
[Schools](#)



1 of 12

Five unit apartment complex within 2 blocks of USC campus, Gate #6. Great for students (most student leases have parents as guarantors). Most USC students live off campus, so housing units like this are always fully leased. Situated on a gated, corner lot, and across from an elementary school, this complex was recently renovated, and has in-unit laundry hook ups, wall-unit AC, and 12 parking spaces. It's within a DPS (Department of Public Safety) and Campus Cruiser patrolled area. This is a great income generating property, not to be missed!

Property Type: Multi-Family  
Community: Downtown Los Angeles  
MLS#: 22176741

Style: Two Level, Low Rise  
County: [Los Angeles](#)

## Property Details for 3620 South BUDLONG, Los Angeles, CA 90007

Details provided by iTech MLS and may not match the public record. [Learn More](#)

### Interior Features

#### Kitchen Information

- Remodeled
- Oven, Range

#### Laundry Information

- Inside Laundry

#### Heating & Cooling

- Wall Cooling Unit(s)

### Multi-Unit Information

#### Community Features

- Units in Complex (Total): 5

#### Multi-Family Information

- # Leased: 5
- # of Buildings: 1
- Owner Pays Water
- Tenant Pays Electricity, Tenant Pays Gas

#### Unit 1 Information

- # of Beds: 2
- # of Baths: 1
- Unfurnished
- Monthly Rent: \$1,700

#### Unit 2 Information

- # of Beds: 3
- # of Baths: 1
- Unfurnished
- Monthly Rent: \$2,250

#### Unit 3 Information

- Unfurnished

#### Unit 4 Information

- # of Beds: 3
- # of Baths: 1
- Unfurnished

#### Unit 5 Information

- Monthly Rent: \$2,350
- # of Beds: 3
- # of Baths: 2
- Unfurnished
- Monthly Rent: \$2,325
- # of Beds: 3
- # of Baths: 1
- Monthly Rent: \$2,250

### Property / Lot Details

#### Property Features

- Automatic Gate, Card/Code Access

- Automatic Gate, Lawn, Sidewalks
- Corner Lot, Near Public Transit

- Tax Parcel Number: 5042017019

#### Lot Information

- Lot Size (Sq. Ft.): 5,548
- Lot Size (Acres): 0.2215
- Lot Size Source: Public Records

#### Property Information

- Updated/Remodeled
- Square Footage Source: Public Records

### Parking / Garage, Exterior Features, Utilities & Financing

#### Parking Information

- # of Parking Spaces (Total): 12
- Parking Space
- Garage

#### Utility Information

- Green Certification Rating: 0.00
- Green Location: Transportation, Walkability
- Green Walk Score: 0
- Green Year Certified: 0

#### Financial Information

- Capitalization Rate (%): 6.25
- Actual Annual Gross Rent: \$128,331
- Gross Rent Multiplier: 11.29

### Building Information

- Total Floors: 2

### Location Details, Misc. Information & Listing Information

#### Location Information

- Cross Streets: W 26th Pl


#### Expense Information

- Operating: \$37,664


#### Listing Information

- Listing Terms: Cash, Cash To Existing Loan
- Buyer Financing: Cash

# Correlation between square footage and sale price




## Roughly linear relationship



$$\text{Sale price} \approx \text{price\_per\_sqft} \times \text{square\_footage} + \text{fixed\_expense}$$

# Data Can be Compactly Represented by Matrices



- Learn parameters ( $w_0, w_1$ ) of the orange line  $y = w_1x + w_0$   
Sq.ft

$$\text{House 1: } 1000 \times w_1 + w_0 = 200,000$$

$$\text{House 2: } 2000 \times w_1 + w_0 = 350,000$$

- Can represent compactly in matrix notation

$$\begin{bmatrix} 1000 & 1 \\ 2000 & 1 \end{bmatrix} \begin{bmatrix} w_1 \\ w_0 \end{bmatrix} = \begin{bmatrix} 200,000 \\ 350,000 \end{bmatrix}$$

# Some Concepts That You Should Know

- Invertibility of Matrices and Computing Inverses
- Vector Norms – L2, Frobenius etc., Inner Products
- Eigenvalues and Eigen-vectors
- Singular Value Decomposition
- Covariance Matrices and Positive Semi-definite-ness

## Excellent Resources:

- Essence of Linear Algebra YouTube Series
- Prof. Gilbert Strang's course at MIT

# Matrix Inverse

- Let us solve the house-price prediction problem

$$\begin{bmatrix} 1000 & 1 \\ 2000 & 1 \end{bmatrix} \begin{bmatrix} w_1 \\ w_0 \end{bmatrix} = \begin{bmatrix} 200,000 \\ 350,000 \end{bmatrix} \quad (1)$$

$$\begin{bmatrix} w_1 \\ w_0 \end{bmatrix} = \left( \begin{bmatrix} 1000 & 1 \\ 2000 & 1 \end{bmatrix} \right)^{-1} \begin{bmatrix} 200,000 \\ 350,000 \end{bmatrix} \quad (2)$$

$$= \frac{1}{-1000} \begin{bmatrix} 1 & -1 \\ -2000 & 1000 \end{bmatrix} \begin{bmatrix} 200,000 \\ 350,000 \end{bmatrix} \quad (3)$$

$$= \frac{1}{-1000} \begin{bmatrix} 150,000 \\ -5 \times 10^7 \end{bmatrix} \quad (4)$$

$$\begin{bmatrix} w_1 \\ w_0 \end{bmatrix} = \begin{bmatrix} 150 \\ 50,000 \end{bmatrix} \quad (5)$$

# You could have data from many houses

- Sale\_price =  
price\_per\_sqft  $\times$  square\_footage + fixed\_expense + unexplainable\_stuff
- Want to learn the price\_per\_sqft and fixed\_expense
- Training data: past sales record.

sqft	sale price
2000	800K
2100	907K
1100	312K
5500	2,600K
...	...

Problem: there isn't a  $\mathbf{w} = [w_1, w_0]^T$  that will satisfy all equations

## Want to predict the best price\_per\_sqft and fixed\_expense

- Sale\_price =  
price\_per\_sqft  $\times$  square\_footage + fixed\_expense + unexplainable\_stuff
- Want to learn the price\_per\_sqft and fixed\_expense
- **Training data:** past sales record.

sqft	sale price	prediction
2000	810K	720K
2100	907K	800K
1100	312K	350K
5500	2,600K	2,600K
...	...	...




# Reduce prediction error

## How to measure errors?

- **absolute** difference:  $|\text{prediction} - \text{sale price}|$ .
- **squared** difference:  $(\text{prediction} - \text{sale price})^2$  [differentiable!].

sqft	sale price	prediction	abs error	squared error
2000	810K	720K	90K	8100
2100	907K	800K	107K	$107^2$
1100	312K	350K	38K	$38^2$
5500	2,600K	2,600K	0	0
...	...			


## Geometric Illustration: Each house corresponds to one line



- Want to find  $\mathbf{w}$  that minimizes the difference between  $\mathbf{X}\mathbf{w}$ ,  $\mathbf{y}$
- But since this a vector, we need an operator that can map the residual vector  $r(\mathbf{w}) = \mathbf{y} - \mathbf{X}\mathbf{w}$  to a scalar

# Norms and Loss Functions

- A vector norm is any function  $f : \mathbb{R}^n \rightarrow \mathbb{R}$  with
  - $f(x) \geq 0$  and  $f(x) = 0 \iff x = 0$
  - $f(ax) = |a|f(x)$  for  $a \in \mathbb{R}$
  - triangle inequality:  $f(x + y) \leq f(x) + f(y)$
- e.g.,  $\ell_2$  norm:  $\|x\|_2 = \sqrt{x^\top x} = \sqrt{\sum_{i=1}^n x_i^2}$
- e.g.,  $\ell_1$  norm:  $\|x\|_1 = \sum_{i=1}^n |x_i|$
- e.g.,  $\ell_\infty$  norm:  $\|x\|_\infty = \max |x_i|$



from inside to outside:  $\ell_1$ ,  $\ell_2$ ,  $\ell_\infty$  norm ball.

# Minimize squared errors

Our model:

Sale\_price =

price\_per\_sqft  $\times$  square\_footage + fixed\_expense + unexplainable\_stuff

Training data:

sqft	sale price	prediction	error	squared error
2000	810K	720K	90K	8100
2100	907K	800K	107K	$107^2$
1100	312K	350K	38K	$38^2$
5500	2,600K	2,600K	0	0
...	...			
Total				$8100 + 107^2 + 38^2 + 0 + \dots$

Aim:

Adjust price\_per\_sqft and fixed\_expense such that the sum of the squared error is minimized — i.e., the unexplainable\_stuff is minimized.

# Linear regression

## Setup:

- **Input:**  $\mathbf{x} \in \mathbb{R}^D$  (covariates, predictors, features, etc)
- **Output:**  $y \in \mathbb{R}$  (responses, targets, outcomes, outputs, etc)
- **Model:**  $f: \mathbf{x} \rightarrow y$ , with  $f(\mathbf{x}) = w_0 + \sum_{d=1}^D w_d x_d = w_0 + \mathbf{w}^\top \mathbf{x}$ .
  - $\mathbf{w} = [w_1 \ w_2 \ \cdots \ w_D]^\top$ : *weights, parameters, or parameter vector*
  - $w_0$  is called *bias*.
  - Sometimes, we also call  $\tilde{\mathbf{w}} = [w_0 \ w_1 \ w_2 \ \cdots \ w_D]^\top$  parameters.
- **Training data:**  $\mathcal{D} = \{(\mathbf{x}_n, y_n), n = 1, 2, \dots, N\}$


## Minimize the Residual sum of squares:

$$RSS(\tilde{\mathbf{w}}) = \sum_{n=1}^N [y_n - f(\mathbf{x}_n)]^2 = \sum_{n=1}^N [y_n - (w_0 + \sum_{d=1}^D w_d x_{nd})]^2$$

## A simple case: $\mathbf{x}$ is just one-dimensional ( $D=1$ )

Residual sum of squares:

$$RSS(\tilde{\mathbf{w}}) = \sum_n [y_n - f(\mathbf{x}_n)]^2 = \sum_n [y_n - (w_0 + w_1 x_n)]^2$$



What kind of function is this? CONVEX (has a unique global minimum)

## A simple case: $\mathbf{x}$ is just one-dimensional ( $D=1$ )

### Residual sum of squares:

$$RSS(\tilde{\mathbf{w}}) = \sum_n [y_n - f(\mathbf{x}_n)]^2 = \sum_n [y_n - (w_0 + w_1 x_n)]^2$$

### Stationary points:

Take derivative with respect to parameters and set it to zero

$$\frac{\partial RSS(\tilde{\mathbf{w}})}{\partial w_0} = 0 \Rightarrow -2 \sum_n [y_n - (w_0 + w_1 x_n)] = 0,$$

$$\frac{\partial RSS(\tilde{\mathbf{w}})}{\partial w_1} = 0 \Rightarrow -2 \sum_n [y_n - (w_0 + w_1 x_n)] x_n = 0.$$

## A simple case: $x$ is just one-dimensional ( $D=1$ )

$$\frac{\partial RSS(\tilde{\mathbf{w}})}{\partial w_0} = 0 \Rightarrow -2 \sum_n [y_n - (w_0 + w_1 x_n)] = 0$$

$$\frac{\partial RSS(\tilde{\mathbf{w}})}{\partial w_1} = 0 \Rightarrow -2 \sum_n [y_n - (w_0 + w_1 x_n)] x_n = 0$$

Simplify these expressions to get the “Normal Equations”:

$$\begin{aligned}\sum y_n &= N w_0 + w_1 \sum x_n \\ \sum x_n y_n &= w_0 \sum x_n + w_1 \sum x_n^2\end{aligned}$$

Solving the system we obtain the **least squares coefficient estimates**:

$$w_1 = \frac{\sum (x_n - \bar{x})(y_n - \bar{y})}{\sum (x_i - \bar{x})^2} \quad \text{and} \quad w_0 = \bar{y} - w_1 \bar{x}$$

where  $\bar{x} = \frac{1}{N} \sum_n x_n$  and  $\bar{y} = \frac{1}{N} \sum_n y_n$ .



## Example

sqft (1000's)	sale price (100k)
1	2
2	3.5
1.5	3
2.5	4.5

### Residual sum of squares:

$$RSS(\tilde{\mathbf{w}}) = \sum_n [y_n - f(\mathbf{x}_n)]^2 = \sum_n [y_n - (w_0 + w_1 x_n)]^2$$

The  $w_1$  and  $w_0$  that minimize this are given by:

$$w_1 = \frac{\sum (x_n - \bar{x})(y_n - \bar{y})}{\sum (x_i - \bar{x})^2} \quad \text{and} \quad w_0 = \bar{y} - w_1 \bar{x}$$

where  $\bar{x} = \frac{1}{N} \sum_n x_n$  and  $\bar{y} = \frac{1}{N} \sum_n y_n$ .

## Example

sqft (1000's)	sale price (100k)
1	2
2	3.5
1.5	3
2.5	4.5

Residual sum of squares:

$$RSS(\tilde{\mathbf{w}}) = \sum_n [y_n - f(\mathbf{x}_n)]^2 = \sum_n [y_n - (w_0 + w_1 x_n)]^2$$

The  $w_1$  and  $w_0$  that minimize this are given by:

$$w_1 \approx 1.6$$

$$w_0 \approx 0.45$$

## Least Mean Squares when $\mathbf{x}$ is $D$ -dimensional

sqft (1000's)	bedrooms	bathrooms	sale price (100k)
1	2	1	2
2	2	2	3.5
1.5	3	2	3
2.5	4	2.5	4.5

**RSS( $\tilde{\mathbf{w}}$ ) in matrix form:**

$$RSS(\tilde{\mathbf{w}}) = \sum_n [y_n - (w_0 + \sum_d w_d x_{nd})]^2 = \sum_n [y_n - \tilde{\mathbf{w}}^\top \tilde{\mathbf{x}}_n]^2,$$

where we have redefined some variables (by augmenting)

$$\tilde{\mathbf{x}} \leftarrow [1 \ x_1 \ x_2 \ \dots \ x_D]^\top, \quad \tilde{\mathbf{w}} \leftarrow [w_0 \ w_1 \ w_2 \ \dots \ w_D]^\top$$

# Least Mean Squares when $\mathbf{x}$ is $D$ -dimensional

$RSS(\tilde{\mathbf{w}})$  in matrix form:

$$RSS(\tilde{\mathbf{w}}) = \sum_n [y_n - (w_0 + \sum_d w_d x_{nd})]^2 = \sum_n [y_n - \tilde{\mathbf{w}}^\top \tilde{\mathbf{x}}_n]^2,$$

where we have redefined some variables (by augmenting)

$$\tilde{\mathbf{x}} \leftarrow [1 \ x_1 \ x_2 \ \dots \ x_D]^\top, \quad \tilde{\mathbf{w}} \leftarrow [w_0 \ w_1 \ w_2 \ \dots \ w_D]^\top$$

which leads to

$$\begin{aligned} RSS(\tilde{\mathbf{w}}) &= \sum_n (y_n - \tilde{\mathbf{w}}^\top \tilde{\mathbf{x}}_n)(y_n - \tilde{\mathbf{x}}_n^\top \tilde{\mathbf{w}}) \\ &= \sum_n \tilde{\mathbf{w}}^\top \tilde{\mathbf{x}}_n \tilde{\mathbf{x}}_n^\top \tilde{\mathbf{w}} - 2y_n \tilde{\mathbf{x}}_n^\top \tilde{\mathbf{w}} + \text{const.} \\ &= \left\{ \tilde{\mathbf{w}}^\top \left( \sum_n \tilde{\mathbf{x}}_n \tilde{\mathbf{x}}_n^\top \right) \tilde{\mathbf{w}} - 2 \left( \sum_n y_n \tilde{\mathbf{x}}_n^\top \right) \tilde{\mathbf{w}} \right\} + \text{const.} \end{aligned}$$

## RSS( $\tilde{\mathbf{w}}$ ) in new notations

From previous slide:

$$RSS(\tilde{\mathbf{w}}) = \left\{ \tilde{\mathbf{w}}^\top \left( \sum_n \tilde{\mathbf{x}}_n \tilde{\mathbf{x}}_n^\top \right) \tilde{\mathbf{w}} - 2 \left( \sum_n y_n \tilde{\mathbf{x}}_n^\top \right) \tilde{\mathbf{w}} \right\} + \text{const.}$$

Design matrix and target vector:

$$\tilde{\mathbf{X}} = \begin{pmatrix} \tilde{\mathbf{x}}_1^\top \\ \tilde{\mathbf{x}}_2^\top \\ \vdots \\ \tilde{\mathbf{x}}_N^\top \end{pmatrix} \in \mathbb{R}^{N \times (D+1)}, \quad \mathbf{y} = \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{pmatrix} \in \mathbb{R}^N$$

Compact expression:

$$RSS(\tilde{\mathbf{w}}) = \|\tilde{\mathbf{X}}\tilde{\mathbf{w}} - \mathbf{y}\|_2^2 = \left\{ \tilde{\mathbf{w}}^\top \tilde{\mathbf{X}}^\top \tilde{\mathbf{X}} \tilde{\mathbf{w}} - 2 \left( \tilde{\mathbf{X}}^\top \mathbf{y} \right)^\top \tilde{\mathbf{w}} \right\} + \text{const}$$

## Example: $RSS(\tilde{\mathbf{w}})$ in compact form

sqft (1000's)	bedrooms	bathrooms	sale price (100k)
1	2	1	2
2	2	2	3.5
1.5	3	2	3
2.5	4	2.5	4.5

Design matrix and target vector:

$$\tilde{\mathbf{X}} = \begin{pmatrix} \tilde{\mathbf{x}}_1^\top \\ \tilde{\mathbf{x}}_2^\top \\ \vdots \\ \tilde{\mathbf{x}}_N^\top \end{pmatrix} \in \mathbb{R}^{N \times (D+1)}, \quad \mathbf{y} = \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{pmatrix} \in \mathbb{R}^N$$

. Compact expression:

$$RSS(\tilde{\mathbf{w}}) = \|\tilde{\mathbf{X}}\tilde{\mathbf{w}} - \mathbf{y}\|_2^2 = \left\{ \tilde{\mathbf{w}}^\top \tilde{\mathbf{X}}^\top \tilde{\mathbf{X}} \tilde{\mathbf{w}} - 2 \left( \tilde{\mathbf{X}}^\top \mathbf{y} \right)^\top \tilde{\mathbf{w}} \right\} + \text{const}$$

## Example: $RSS(\tilde{\mathbf{w}})$ in compact form

sqft (1000's)	bedrooms	bathrooms	sale price (100k)
1	2	1	2
2	2	2	3.5
1.5	3	2	3
2.5	4	2.5	4.5

Design matrix and target vector:

$$\tilde{\mathbf{X}} = \begin{pmatrix} \tilde{\mathbf{x}}_1^\top \\ \tilde{\mathbf{x}}_2^\top \\ \vdots \\ \tilde{\mathbf{x}}_N^\top \end{pmatrix} = \begin{bmatrix} 1 & 1 & 2 & 1 \\ 1 & 2 & 2 & 2 \\ 1 & 1.5 & 3 & 2 \\ 1 & 2.5 & 4 & 2.5 \end{bmatrix}, \quad \mathbf{y} = \begin{bmatrix} 2 \\ 3.5 \\ 3 \\ 4.5 \end{bmatrix}$$

. Compact expression:

$$RSS(\tilde{\mathbf{w}}) = \|\tilde{\mathbf{X}}\tilde{\mathbf{w}} - \mathbf{y}\|_2^2 = \left\{ \tilde{\mathbf{w}}^\top \tilde{\mathbf{X}}^\top \tilde{\mathbf{X}} \tilde{\mathbf{w}} - 2 \left( \tilde{\mathbf{X}}^\top \mathbf{y} \right)^\top \tilde{\mathbf{w}} \right\} + \text{const}$$

# Solution in matrix form

## Compact expression

$$RSS(\tilde{\mathbf{w}}) = ||\tilde{\mathbf{X}}\tilde{\mathbf{w}} - \mathbf{y}||_2^2 = \left\{ \tilde{\mathbf{w}}^\top \tilde{\mathbf{X}}^\top \tilde{\mathbf{X}} \tilde{\mathbf{w}} - 2 \left( \tilde{\mathbf{X}}^\top \mathbf{y} \right)^\top \tilde{\mathbf{w}} \right\} + \text{const}$$

## Gradients of Linear and Quadratic Functions

- $\nabla_{\mathbf{x}}(\mathbf{b}^\top \mathbf{x}) = \mathbf{b}$
- $\nabla_{\mathbf{x}}(\mathbf{x}^\top \mathbf{A} \mathbf{x}) = 2\mathbf{A} \mathbf{x}$  (symmetric  $\mathbf{A}$ )

## Normal equation

$$\nabla_{\tilde{\mathbf{w}}} RSS(\tilde{\mathbf{w}}) = 2\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}} \tilde{\mathbf{w}} - 2\tilde{\mathbf{X}}^\top \mathbf{y} = 0$$

This leads to the **least-mean-squares** (LMS) solution

$$\tilde{\mathbf{w}}^{LMS} = \left( \tilde{\mathbf{X}}^\top \tilde{\mathbf{X}} \right)^{-1} \tilde{\mathbf{X}}^\top \mathbf{y}$$



## Example: $RSS(\tilde{\mathbf{w}})$ in compact form

sqft (1000's)	bedrooms	bathrooms	sale price (100k)
1	2	1	2
2	2	2	3.5
1.5	3	2	3
2.5	4	2.5	4.5

Write the **least-mean-squares** (LMS) solution

$$\tilde{\mathbf{w}}^{LMS} = \left( \tilde{\mathbf{X}}^T \tilde{\mathbf{X}} \right)^{-1} \tilde{\mathbf{X}}^T \mathbf{y}$$

Can use solvers in Matlab, Python etc., to compute this for any given  $\tilde{\mathbf{X}}$  and  $\mathbf{y}$ .

## Exercise: $RSS(\tilde{\mathbf{w}})$ in compact form

Using the general **least-mean-squares** (LMS) solution

$$\tilde{\mathbf{w}}^{LMS} = \left( \tilde{\mathbf{X}}^\top \tilde{\mathbf{X}} \right)^{-1} \tilde{\mathbf{X}}^\top \mathbf{y}$$

recover the uni-variate solution that we had computed earlier:

$$w_1 = \frac{\sum (x_n - \bar{x})(y_n - \bar{y})}{\sum (x_i - \bar{x})^2} \quad \text{and} \quad w_0 = \bar{y} - w_1 \bar{x}$$

where  $\bar{x} = \frac{1}{N} \sum_n x_n$  and  $\bar{y} = \frac{1}{N} \sum_n y_n$ .


## Exercise: $RSS(\tilde{\mathbf{w}})$ in compact form

For the 1-D case, the **least-mean-squares** solution is

$$\begin{aligned}\tilde{\mathbf{w}}^{LMS} &= (\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}})^{-1} \tilde{\mathbf{X}}^\top \mathbf{y} \\&= \left( \begin{bmatrix} 1 & 1 & \dots & 1 \\ x_1 & x_2 & \dots & x_N \end{bmatrix} \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ 1 & \dots \\ 1 & x_N \end{bmatrix} \right)^{-1} \begin{bmatrix} 1 & 1 & \dots & 1 \\ x_1 & x_2 & \dots & x_N \end{bmatrix} \begin{bmatrix} y_1 \\ y_2 \\ \dots \\ y_N \end{bmatrix} \\&= \left( \begin{bmatrix} N & N\bar{x} \\ N\bar{x} & \sum_n x_n^2 \end{bmatrix} \right)^{-1} \begin{bmatrix} \sum_n y_n \\ \sum_n x_n y_n \end{bmatrix} \\ \begin{bmatrix} w_0 \\ w_1 \end{bmatrix} &= \frac{1}{\sum (x_i - \bar{x})^2} \begin{bmatrix} \bar{y} \sum (x_i - \bar{x})^2 - \bar{x} \sum (x_n - \bar{x})(y_n - \bar{y}) \\ \sum (x_n - \bar{x})(y_n - \bar{y}) \end{bmatrix}\end{aligned}$$

where  $\bar{x} = \frac{1}{N} \sum_n x_n$  and  $\bar{y} = \frac{1}{N} \sum_n y_n$ .

# Why is minimizing RSS sensible?



- Want to find  $\mathbf{w}$  that minimizes the difference between  $\mathbf{X}\mathbf{w}$ ,  $\mathbf{y}$
- But since this a vector, we need an operator that can map the residual vector  $r(\mathbf{w}) = \mathbf{y} - \mathbf{X}\mathbf{w}$  to a scalar
- We take the sum of the squares of the elements of  $r(\mathbf{w})$

# Why is minimizing RSS sensible?

## Probabilistic interpretation

- Noisy observation model:

$$Y = w_0 + w_1 X + \eta$$

where  $\eta \sim N(0, \sigma^2)$  is a Gaussian random variable

- Conditional likelihood of one training sample:

$$p(y_n|x_n) = N(w_0 + w_1 x_n, \sigma^2) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{[y_n - (w_0 + w_1 x_n)]^2}{2\sigma^2}}$$

## Probabilistic interpretation (cont'd)

Log-likelihood of the training data  $\mathcal{D}$  (assuming i.i.d):

$$\begin{aligned}\log P(\mathcal{D}) &= \log \prod_{n=1}^N p(y_n|x_n) = \sum_n \log p(y_n|x_n) \\&= \sum_n \left\{ -\frac{[y_n - (w_0 + w_1 x_n)]^2}{2\sigma^2} - \log \sqrt{2\pi}\sigma \right\} \\&= -\frac{1}{2\sigma^2} \sum_n [y_n - (w_0 + w_1 x_n)]^2 - \frac{N}{2} \log \sigma^2 - N \log \sqrt{2\pi} \\&= -\frac{1}{2} \left\{ \frac{1}{\sigma^2} \sum_n [y_n - (w_0 + w_1 x_n)]^2 + N \log \sigma^2 \right\} + \text{const}\end{aligned}$$

What is the relationship between minimizing RSS and maximizing the log-likelihood?

# Maximum likelihood estimation

Estimating  $\sigma$ ,  $w_0$  and  $w_1$  can be done in two steps

- Maximize over  $w_0$  and  $w_1$ :

$$\max \log P(\mathcal{D}) \Leftrightarrow \min \sum_n [y_n - (w_0 + w_1 x_n)]^2 \leftarrow \text{This is RSS}(\tilde{\mathbf{w}})!$$

- Maximize over  $s = \sigma^2$ :

$$\begin{aligned} \frac{\partial \log P(\mathcal{D})}{\partial s} &= -\frac{1}{2} \left\{ -\frac{1}{s^2} \sum_n [y_n - (w_0 + w_1 x_n)]^2 + N \frac{1}{s} \right\} = 0 \\ \rightarrow \sigma^{*2} = s^* &= \frac{1}{N} \sum_n [y_n - (w_0 + w_1 x_n)]^2 \end{aligned}$$

## How does this probabilistic interpretation help us?

- It gives a solid footing to our intuition: minimizing  $\text{RSS}(\tilde{\mathbf{w}})$  is a sensible thing based on reasonable modeling assumptions.
- Estimating  $\sigma^*$  tells us how much noise there is in our predictions. For example, it allows us to place confidence intervals around our predictions.



# Computational complexity of the Least Squares Solution

Bottleneck of computing the solution?

$$\mathbf{w} = \left( \mathbf{X}^\top \mathbf{X} \right)^{-1} \mathbf{X} \mathbf{y}$$

Matrix multiply of  $\mathbf{X}^\top \mathbf{X} \in \mathbb{R}^{(D+1) \times (D+1)}$

Inverting the matrix  $\mathbf{X}^\top \mathbf{X}$


How many operations do we need?

- $O(ND^2)$  for matrix multiplication
- $O(D^3)$  (e.g., using Gauss-Jordan elimination) or  $O(D^{2.373})$  (recent theoretical advances) for matrix inversion
- Impractical for very large  $D$  or  $N$

# Alternative method: Batch Gradient Descent

## (Batch) Gradient descent

- Initialize  $\mathbf{w}$  to  $\mathbf{w}^{(0)}$  (e.g., randomly);  
set  $t = 0$ ; choose  $\eta > 0$
- Loop *until convergence*
  1. Compute the gradient
$$\nabla \text{RSS}(\mathbf{w}) = \mathbf{X}^\top (\mathbf{X}\mathbf{w}^{(t)} - \mathbf{y})$$
  2. Update the parameters
$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \eta \nabla \text{RSS}(\mathbf{w})$$
  3.  $t \leftarrow t + 1$



What is the complexity of each iteration?

$O(\text{ND})$

# Why would this work?

If gradient descent converges, it will converge to the same solution as using matrix inversion.

This is because  $RSS(\mathbf{w})$  is a convex function in its parameters  $\mathbf{w}$

Hessian of RSS

$$\begin{aligned}RSS(\mathbf{w}) &= \mathbf{w}^\top \mathbf{X}^\top \mathbf{X} \mathbf{w} - 2 (\mathbf{X}^\top \mathbf{y})^\top \mathbf{w} + \text{const} \\ \Rightarrow \frac{\partial^2 RSS(\mathbf{w})}{\partial \mathbf{w} \mathbf{w}^\top} &= 2 \mathbf{X}^\top \mathbf{X}\end{aligned}$$


$\mathbf{X}^\top \mathbf{X}$  is positive semidefinite, because for any  $\mathbf{v}$

$$\mathbf{v}^\top \mathbf{X}^\top \mathbf{X} \mathbf{v} = \|\mathbf{X}^\top \mathbf{v}\|_2^2 \geq 0$$

# Alternative method: Batch Gradient Descent

## (Batch) Gradient descent

- Initialize  $\mathbf{w}$  to  $\mathbf{w}^{(0)}$  (e.g., randomly);  
set  $t = 0$ ; choose  $\eta > 0$
- Loop *until convergence*
  1. Compute the gradient
$$\nabla \text{RSS}(\mathbf{w}) = \mathbf{X}^\top \mathbf{X} \mathbf{w}^{(t)} - \mathbf{X}^\top \mathbf{y}$$
  2. Update the parameters
$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \eta \nabla \text{RSS}(\mathbf{w})$$
  3.  $t \leftarrow t + 1$



What is the complexity of each iteration?

$O(\text{ND})$

# Stochastic gradient descent (SGD)

**Widrow-Hoff rule:** update parameters using one example at a time

- Initialize  $\mathbf{w}$  to some  $\mathbf{w}^{(0)}$ ; set  $t = 0$ ; choose  $\eta > 0$
- Loop *until convergence*
  1. random choose a training a sample  $\mathbf{x}_t$
  2. Compute its contribution to the gradient

$$\mathbf{g}_t = (\mathbf{x}_t^\top \mathbf{w}^{(t)} - y_t) \mathbf{x}_t$$


3. Update the parameters
$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \eta \mathbf{g}_t$$
4.  $t \leftarrow t + 1$

How does the complexity per iteration compare with gradient descent?


- $O(ND)$  for gradient descent versus  $O(D)$  for SGD

# SGD versus Batch GD

Stochastic Gradient Descent




Gradient Descent



- SGD reduces per-iteration complexity from  $O(ND)$  to  $O(D)$
- But it is noisier and can take longer to converge

# How to Choose Learning Rate $\eta$ in practice?

- Try 0.0001, 0.001, 0.01, 0.1 etc. on a validation dataset (more on this later) and choose the one that gives fastest, stable convergence
- Reduce  $\eta$  by a constant factor (eg. 10) when learning saturates so that we can reach closer to the true minimum.
- More advanced learning rate schedules such as AdaGrad, Adam, AdaDelta are used in practice.



# Mini-Summary

- Linear regression is the linear combination of features  
 $f : \mathbf{x} \rightarrow y$ , with  $f(\mathbf{x}) = w_0 + \sum_d w_d x_d = w_0 + \mathbf{w}^\top \mathbf{x}$
- If we minimize residual sum of squares as our learning objective, we get a closed-form solution of parameters
- Probabilistic interpretation: maximum likelihood if assuming residual is Gaussian distributed
- Gradient Descent and mini-batch SGD can overcome computational issues